

Comparison of overall ratings of Airbnb to Pipeline

June 21, 2016

```
In [1]: %matplotlib inline
```

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from __future__ import division
from pandas.tools.plotting import autocorrelation_plot
from scipy.stats import ks_2samp
```

```
pd.set_option('display.mpl_style', 'default') # Make the graphs a bit prettier
plt.rcParams['figure.figsize'] = (15, 5)
plt.rcParams['font.family'] = 'sans-serif'
```

```
# This is necessary to show lots of columns in pandas 0.12.
# Not necessary in pandas 0.13.
pd.set_option('display.width', 5000)
pd.set_option('display.max_columns', 60)
```

```
c:\python27\lib\site-packages\IPython\core\interactiveshell.py:2885: FutureWarning:
mpl_style had been deprecated and will be removed in a future version.
Use `matplotlib.pyplot.style.use` instead.
```

```
exec(code_obj, self.user_global_ns, self.user_ns)
```

```
In [2]: pipeline = pd.read_csv('C:/Python27/output_improved_AMS.csv')[['Listing ID', 'Overall Rating']]
airbnb = pd.read_csv('C:/Python27/AirbnbRating.csv')[['Id', 'Overall']]
```

```
In [3]: # Get the overall ratings for each listing from the pipeline and save
# them in a new csv file
```

```
set1=pipeline.groupby('Listing ID').mean().round(2)
set1.to_csv(path_or_buf='C:/Python27/Pipelineclean1.csv')
```

```
In [4]: # Delete all the rows with a NaN value and save in a new csv. Both the csv
# joined using the command line
# and based on Listing ID
```

```
set2 = airbnb.dropna()
set2.to_csv(path_or_buf='C:/Python27/Airbnbclean.csv')
```

```
In [5]: # Read the csv of the merged file
```

```
comparison=pd.read_csv('C:/Python27/merged_AMS.csv')
comparison[:3]
```

```
Out[5]:
```

	Airbnb	Listing ID	Pipeline
0	5.0	1000126	4.0
1	5.0	1000252	4.0
2	5.0	1000866	4.5

0.1 Comparison of Airbnb star rating and pipeline

For every listing the stars given by pipeline are compared to the stars of Airbnb. Thus, the two values form a combination. All the possible combinations are presented in the table

```
In [7]: # For filling the dataframe
```

```
i=5.0
s0=comparison[comparison['Airbnb']==i]
s00=s0['Pipeline'].value_counts()
b=pd.DataFrame({'Airbnb':i, 'Pipeline':s00.index,
                'V_Frequency': s00.values})

i=4.5
s1=comparison[comparison['Airbnb']==i]
s11=s1['Pipeline'].value_counts()
b1=pd.DataFrame({'Airbnb':i, 'Pipeline':s11.index,
                'V_Frequency': s11.values})

i=4.0
s2=comparison[comparison['Airbnb']==i]
s22=s2['Pipeline'].value_counts()
b2=pd.DataFrame({'Airbnb':i, 'Pipeline':s22.index,
                'V_Frequency': s22.values})

i=3.5
s3=comparison[comparison['Airbnb']==i]
s33=s3['Pipeline'].value_counts()
b3=pd.DataFrame({'Airbnb':i, 'Pipeline':s33.index,
                'V_Frequency': s33.values})

i=3.0
s4=comparison[comparison['Airbnb']==i]
s44=s4['Pipeline'].value_counts()
b4=pd.DataFrame({'Airbnb':i, 'Pipeline':s44.index,
                'V_Frequency': s44.values},index=[12])

stars_compared=pd.concat([b,b1,b2,b3,b4],ignore_index=True)
stars_compared['Percentage']=
((stars_compared['V_Frequency']/stars_compared['V_Frequency'].sum())*100)
```

```
stars_compared.columns=['Airbnb','Pipeline','Frequency','Percentage']
stars_compared
```

```
Out [7]:
```

	Airbnb	Pipeline	Frequency	Percentage
0	5.0	4.0	650	31.355523
1	5.0	4.5	404	19.488664
2	5.0	3.5	2	0.096479
3	4.5	4.0	686	33.092137
4	4.5	4.5	199	9.599614
5	4.5	3.5	4	0.192957
6	4.5	3.0	1	0.048239
7	4.0	4.0	100	4.823927
8	4.0	4.5	11	0.530632
9	4.0	3.5	7	0.337675
10	3.5	4.0	6	0.289436
11	3.5	3.5	2	0.096479
12	3.0	5.0	1	0.048239

0.2 Visualization of most frequent star values and grouping non-frequent as “others”

Since we have combinations which appear only once or twice in the whole corpus, then we group these into “other”. The pie-chart shows the most frequent combinations.

```
In [8]: ot=stars_compared[stars_compared['Percentage']<0.55]
other=ot['Percentage'].sum()
other
freq=ot['Frequency'].sum()
main=stars_compared[stars_compared['Percentage']>=0.55]
main.loc[13]=['Other','Other',freq,other]
main
```

```
c:\python27\lib\site-packages\ipykernel\__main__.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```


See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/>

```
Out [8]:
```

	Airbnb	Pipeline	Frequency	Percentage
0	5	4	650	31.355523
1	5	4.5	404	19.488664
3	4.5	4	686	33.092137
4	4.5	4.5	199	9.599614
7	4	4	100	4.823927
13	Other	Other	34	1.640135

```
In [9]: main['Percentage'].plot(kind='pie', labels=['(5 : 4)', '(5 : 4.5)',
          '(4.5 : 4)', '(4.5 : 4.5)', '(4 : 4)', 'Others'], autopct='%.2f',
          fontsize=10, figsize=(6, 6))
```

Out [9]: <matplotlib.axes._subplots.AxesSubplot at 0xd23f7b0>



0.3 The occurrence of “rare” combinations


And here we have the occurrence of non-frequent combinations, all less than 0.55%. Below they are also visualized in a bar chart with their corresponding frequencies.

In [10]: ot

Out [10]:

	Airbnb	Pipeline	Frequency	Percentage
2	5.0	3.5	2	0.096479
5	4.5	3.5	4	0.192957
6	4.5	3.0	1	0.048239
8	4.0	4.5	11	0.530632
9	4.0	3.5	7	0.337675
10	3.5	4.0	6	0.289436
11	3.5	3.5	2	0.096479
12	3.0	5.0	1	0.048239

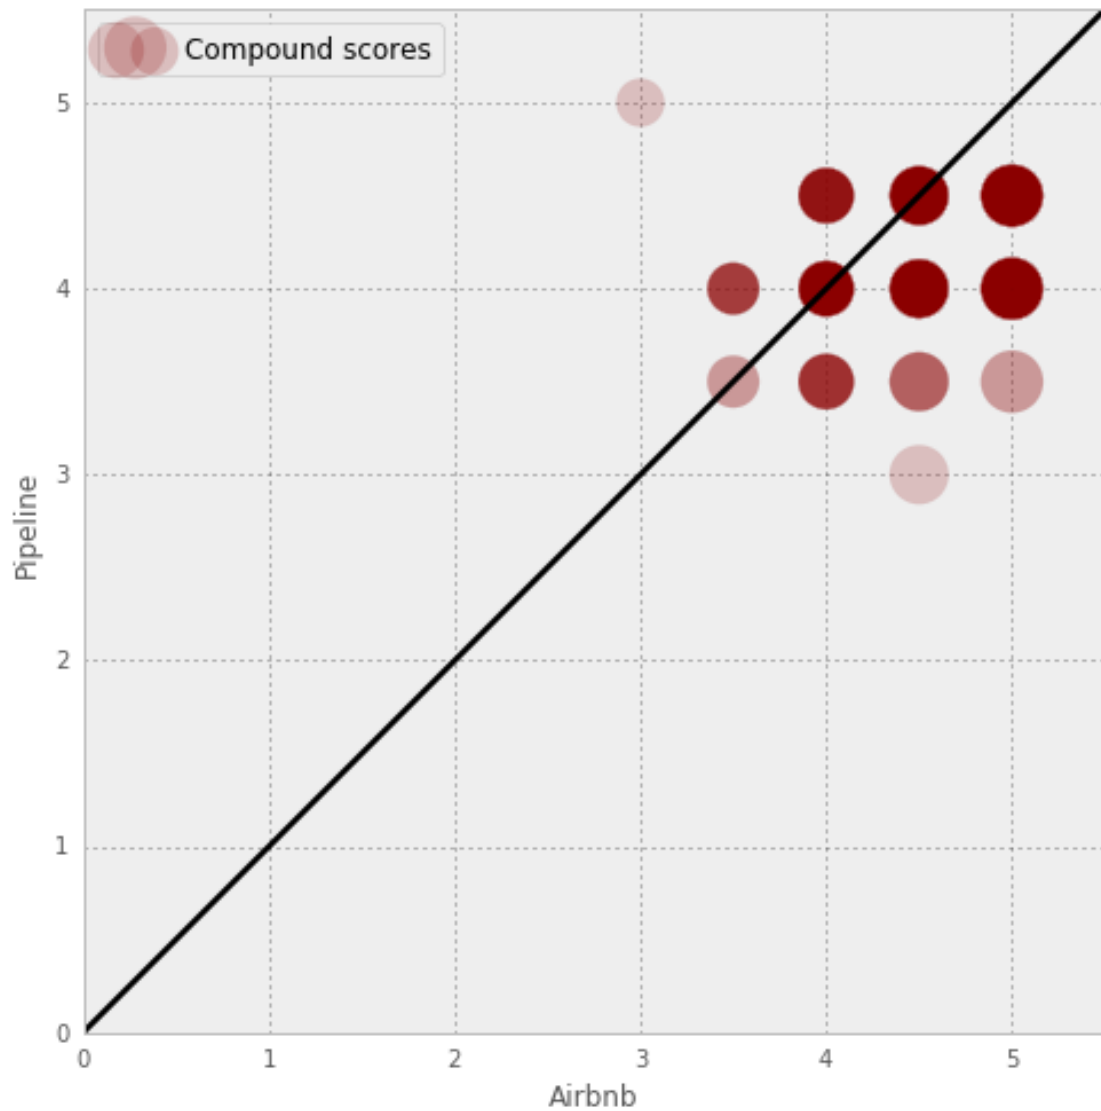
```
In [11]: ot['Percentage'].plot(kind='bar',figsize=(6,4))
plt.axis([-1, 14, 0, 0.7])
plt.show()
```



0.4 How value combination of stars is distributed

The scatter shows how the combinations of stars are distributed. We see only the main outliers (3.0,5.0), (4.5,3.0) and (5.0, 3.5)

```
In [12]: fx = comparison.plot(kind='scatter', x='Airbnb', y='Pipeline', color='Dark
label='Compound scores', s=comparison['Airbnb']*150,
alpha=0.2, figsize=(8,8)).set_xlim(0,8)
line = plt.plot([0,1,2,3,4,5,6], [0,1,2,3,4,5,6])
plt.axis([0, 5.5, 0, 5.5])
plt.setp(line, color='Black', linewidth=2.5)
plt.xlabel('Airbnb')
plt.ylabel('Pipeline')
plt.show()
```



0.5 The frequency for all Airbnb assigned stars

We see that for Airbnb the most assigned value is 5.0 with 50.95 % and also 4.5 with 42.93%.


```
In [13]: abnb=stars_compared[['Airbnb', 'Frequency', 'Percentage']].groupby('Airbnb')
          abnb
```

```
Out[13]:
```

Airbnb	Frequency	Percentage
3.0	1	0.048239
3.5	8	0.385914
4.0	118	5.692233
4.5	890	42.932947
5.0	1056	50.940666

```
In [14]: abnb['Percentage'].plot(kind='bar',figsize=(5,5))
plt.axis([0, 5.5, 0, 80])
```

```
Out[14]: [0, 5.5, 0, 80]
```



0.6 The frequency for the stars assigned by the pipeline

for the pipeline things change a bit, as we have the most assigned value 4.0, followed by 4.5


```
In [15]: pip=stars_compared[['Pipeline','Frequency','Percentage']].groupby('Pipeline')
pip
```

```
Out[15]:
```

Pipeline	Frequency	Percentage
3.0	1	0.048239
3.5	15	0.723589
4.0	1442	69.561023
4.5	614	29.618910
5.0	1	0.048239

```
In [16]: pip['Percentage'].plot(kind='bar',color='DarkRed',figsize=(5,5))
plt.axis([0, 5.5, 0, 80])
```

```
Out[16]: [0, 5.5, 0, 80]
```



0.7 The difference in stars Airbnb and pipeline compared

The difference in stars for every listing is calculated and the frequency of differences is shown in the table. We notice that the highest difference is -2.0, followed by three cases with difference 1.5

```
In [17]: comparison['Difference']=comparison['Airbnb']-comparison['Pipeline']
dfr=comparison['Difference'].value_counts()
norm= comparison['Difference'].value_counts(normalize=True)
dfrf=dfr.to_frame()
dfrf['Normalized']=norm
dfrf.columns=['Frequency','Normalized']
dfrf
```

```
Out[17]:
```

	Frequency	Normalized
0.5	1097	0.529185

1.0	654	0.315485
0.0	301	0.145200
-0.5	17	0.008201
1.5	3	0.001447
-2.0	1	0.000482

0.8 Get Listing ID for cases with big differences

We identify the listings with big differences in order to get what is wrong with them.

```
In [18]: bg1=comparison[comparison['Difference']>1]
ls1=comparison[comparison['Difference']<-1]
weird=pd.concat([bg1,ls1])
weird_listing=weird['Listing ID']
weird
```

```
Out[18]:
```

	Airbnb	Listing ID	Pipeline	Difference
491	5.0	1357971	3.5	1.5
564	4.5	1410370	3.0	1.5
1807	5.0	2606699	3.5	1.5
34	3.0	1022631	5.0	-2.0

0.9 Let's check these cases

Starting from the one with the biggest difference. So only one review... Which has a good sentiment score and makes the whole rating of pipeline base only on this review

```
In [19]: full_content = pd.read_csv('C:/Python27/output_improved_AMS.csv')
full_content[full_content['Listing ID']==1022631]
```

```
Out[19]:
```

	Listing ID	Reviewer ID	Review ID	
6515	1022631	9191748	10747505	Had a lovely break, very cosy ho

1 And now the cases with difference 1.5

1.1 Let's see this one

And this other place has no reviews except of one cancelled reservation, for which the pipeline is neutral (3 stars)

```
In [20]: full_content[full_content['Listing ID']==1410370]
```

```
Out[20]:
```

	Listing ID	Reviewer ID	Review ID	
90721	1410370	4091313	30601951	The reservation was canceled 20
90722	1410370	4091313	30601951	This is an

1.2 And another ...

```
In [21]: full_content[full_content['Listing ID']==1357971]
```

```
Out[21]:
```

	Listing ID	Reviewer ID	Review ID	
79596	1357971	8022631	7741708	We had a great time in Josh's C
79597	1357971	8022631	7741708	So I hope
79598	1357971	8022631	7741708	Thanks for the good ti
79599	1357971	8545449	8548435	A dream came
79600	1357971	8545449	8548435	Its a great experience to be li
79601	1357971	8545449	8548435	Almost like forgot
79602	1357971	8545449	8548435	I high
79603	1357971	8545449	8548435	
79604	1357971	9365972	9053128	We had a lovely stay in Joshua
79605	1357971	13033115	31327563	Josh is a very flexible, friendl
79606	1357971	13033115	31327563	Angela and Josh are very ver
79607	1357971	13033115	31327563	unfortunately we could not stay
79608	1357971	13033115	31327563	always happy :) next time it w
79609	1357971	13033115	31327563	Greetings f
79610	1357971	34536732	37758257	An excellent choice for an inex
79611	1357971	34536732	37758257	This is not luxury living but t
79612	1357971	39033958	39158959	The host canceled this reservat
79613	1357971	39033958	39158959	This is an
79614	1357971	14709263	53183435	I am sorry to say my friend and
79615	1357971	14709263	53183435	What you see on the pictures is
79616	1357971	14709263	53183435	What you get is much more m
79617	1357971	14709263	53183435	When Joshua was introducing us
79618	1357971	14709263	53183435	and above all a terribly bad sm
79619	1357971	14709263	53183435	The sink was covered with coffe
79620	1357971	14709263	53183435	Thus water was stagnating
79621	1357971	14709263	53183435	There was a crazy mess in the f
79622	1357971	14709263	53183435	The heater that you get creates
79623	1357971	14709263	53183435	I really cannot advise this can
79624	1357971	14709263	53183435	It seems like Josh wants to mak
79625	1357971	14709263	53183435	Talking about honesty, the peop
79626	1357971	14709263	53183435	That's maybe why Joshua insiste
79627	1357971	14709263	53183435	So it seems like Josh does not
79628	1357971	14709263	53183435	I will add that after we did th
79629	1357971	14709263	53183435	Seems like once you've payed, y
79630	1357971	14709263	53183435	I am truly sorry for this bad r
79631	1357971	14709263	53183435	I still want to say that the be

1.3 And the last case

```
In [22]: full_content[full_content['Listing ID']==2606699]
```

```
Out[22]:
```

	Listing ID	Reviewer ID	Review ID	
268229	2606699	738418	17075400	Amazing home ar
268230	2606699	738418	17075400	Rob was very frie

268231	2606699	738418	17075400	Highly recom
268232	2606699	16610752	20933169	I stayed at Rob's house over t
268233	2606699	16610752	20933169	The house is
268234	2606699	16610752	20933169	It's in a quiet neighborhood b
268235	2606699	16610752	20933169	Rob was very nice and accommo
268236	2606699	16610752	20933169	I would gladl
268237	2606699	21004608	28894069	the location was amazing, the
268238	2606699	21004608	28894069	however there was one bedroom
268239	2606699	32418854	43326085	It was a great home base while
268240	2606699	32418854	43326085	There are lots of stairs, but
268241	2606699	32418854	43326085	The apartment was exa
268242	2606699	32418854	43326085	Rob was
268243	2606699	32418854	43326085	Rob was there at our sched
268244	2606699	32418854	43326085	Comfy be
268245	2606699	32418854	43326085	W

1.4 RMSE of differences in two samples

```
In [23]: SSD=comparison['Difference'].pow(2).sum()
n=comparison['Difference'].count()
MSE=SSD/n
RMSE=np.sqrt(MSE)
print RMSE.round(3)
```

0.675

```
In [24]: # The common difference between pipeline and Airbnb
bs=comparison['Difference'].mean()
bs
```

Out [24]: 0.57718282682103228

```
In [25]: # STANDARD DEVIATION
comparison['Difference'].std()
```

Out [25]: 0.34919291093687588

1.5 Let's now remove half a star from Airbnb values

We saw that the difference of half a star in Airbnb values higher than the pipeline prevails. Therefore known that there is a bias towards high scores in the Airbnb system, we try to subtract from all the values half a star.

```
In [26]: comparison['RE: Airbnb']=comparison['Airbnb']-0.5
comparison['RE: Difference']=comparison['RE: Airbnb'] - comparison['Pipeline']
```

```
In [27]: dfr1=comparison['RE: Difference'].value_counts()
norm1= comparison['RE: Difference'].value_counts(normalize=True)
```

```

dfrf1=dfrf1.to_frame()
dfrf1 ['Frequency'] = dfrf['Frequency']
dfrf1['RE: Normalized']=norm1
dfrf1['Normalized']=norm
dfrf1.columns=['RE: Frequency','Frequency','RE: Normalized','Normalized']
dfrf1

```

```

Out [27]:
      RE: Frequency  Frequency  RE: Normalized  Normalized
0.0              1097      301.0         0.529185    0.145200
0.5               654     1097.0         0.315485    0.529185
-0.5              301       17.0         0.145200    0.008201
-1.0               17        NaN         0.008201         NaN
1.0                3      654.0         0.001447    0.315485
-2.5               1        NaN         0.000482         NaN

```

1.6 RMSE in the new set

We see that the RMSE is reduced with 47%.

```

In [28]: SSD1=comparison['RE: Difference'].pow(2).sum()
MSE_RE=SSD1/n
RMSE_RE=np.sqrt(MSE_RE)
RMSE_RE

```

```

Out [28]: 0.35753888868532163

```

```

In [31]: (RMSE_RE-RMSE)/RMSE

```

```

Out [31]: -0.46995892022441427

```

1.7 Get the IDs of Listings with less than 3 reviews

Another test is the case of listings with less than 3 reviews, as we said that their sentiment value would not be very reliable. We see that in 213 listings we have less than 3 reviews.

```

In [32]: a=full_content[['Listing ID','Review ID']]
a_nodup=a.drop_duplicates()
count=a_nodup.groupby('Listing ID').count()
count['Listing ID']=count.index
d = count[count['Review ID']<3]
less_3 = d[['Listing ID']]
less_3.count()

```

```

Out [32]: Listing ID      213
dtype: int64

```

1.8 Remove from analysis of differences these cases

For having reliable scores we remove these cases from the analysis and then we check again the differences. We see that in the new case the maximal difference will be 1.5 stars in only 2 case. So we have omitted the case with -2 stars difference and 1 case with 1.5 stars difference. The result means that listings with a high number of reviews generate more reliable scores, however since we have in most of the cases more than 3 reviews per listing, it would not directly affect the pipeline.

```
In [33]: mutual=comparison[comparison['Listing ID'].isin(less_3['Listing ID'])]
         indexes_ID=mutual.index
         for i in less_3['Listing ID']:
             new_comparison=comparison.drop(indexes_ID)
         new_comparison[:5]
```

```
Out [33]:
```

	Airbnb	Listing ID	Pipeline	Difference	RE: Airbnb	RE: Difference
0	5.0	1000126	4.0	1.0	4.5	0.5
1	5.0	1000252	4.0	1.0	4.5	0.5
2	5.0	1000866	4.5	0.5	4.5	0.0
3	4.5	1001885	4.5	0.0	4.0	-0.5
4	4.5	1002180	4.5	0.0	4.0	-0.5

```
In [34]: new_comparison['Difference']=new_comparison['Airbnb']-new_comparison['Pipeline']
         dfr_new=new_comparison['Difference'].value_counts()
         norm_new = new_comparison['Difference'].value_counts(normalize=True)
         dfrf_new=dfr_new.to_frame()
         dfrf_new['Normalized']=norm_new
         dfrf_new.columns=['New_Frequency', 'New_Normalized']
         dfrf_new
```

```
Out [34]:
```

	New_Frequency	New_Normalized
0.5	1081	0.530422
1.0	644	0.315996
0.0	294	0.144259
-0.5	17	0.008342
1.5	2	0.000981

1.9 RMSE of the set of listings with 3+ reviews

After calculating the new RSME we see that indeed there is no significant change because the cases with less than 3 reviews are very few.

```
In [36]: SSD3=new_comparison['RE: Difference'].pow(2).sum()
         MSE_RE_new=SSD3/new_comparison['RE: Difference'].count()
         RMSE_RE_new=np.sqrt(MSE_RE_new)
         RMSE_RE_new
```

```
Out [36]: 0.35268492111506988
```

```
In [37]: RMSE_RE_new-RMSE_RE
```

```
Out [37]: -0.0048539675702517493
```