
Feature-based opinion mining: A proposal to enhance customer focus based on text feedback



UNIVERSITEIT VAN AMSTERDAM

Business Information Systems
MSc. INFORMATION STUDIES

Author: Antigoni Kourou
Supervisor: Dr. Maarten Marx

Student number: 11118172
Signature:

JUNE 2016

DEDICATION

Here goes the dedication.

AUTHOR'S DECLARATION

I declare that the work in this dissertation was carried out in accordance with the requirements of the University's Regulations and Code of Practice for Research Degree Programmes and that it has not been submitted for any other academic award. Except where indicated by specific reference in the text, the work is the candidate's own work. Work done in collaboration with, or with the assistance of, others, is indicated as such. Any views expressed in the dissertation are those of the author.

SIGNED: DATE:

ABSTRACT

The increased use of Internet and Information and Communication Technologies (ICTs) have affected radically the tourism industry. Current online travel systems, aiming to assist the consumer in finding the suitable offers, provide a great complexity and diversity in offers to fulfill their customers' needs. However, it is incredibly difficult for service providers to understand their customers' needs and preferences. ICTs and Internet have transformed the markets to customer-driven [?], meaning that the users are co-creators of traveling information through blogs, forums, social media and feedback systems. Customer-created content is found mainly in the form of free-text opinions, which require a tremendous amount of work to be manually analyzed and get the proper insights. The tools for automating this analysis fall into the category of Natural Language Processing (NLP), explicitly sentiment analysis. This paper proposes the use of feature-based sentiment mining for estimating discrete score for each accommodation feature, based on reviews of feedback system. The proposed approach comes in three stages: (i) feature extraction from full text reviews based on accommodation ontology; (ii) sentiment detection for each feature on sentence-level and hybrid-based method, including both rule-based and lexicon approaches; and (iii) data analysis of discrete sentiment scores per feature and overall. The aim of this proposal is to show how insights from text content can be aligned with the quantitative data of feedback systems for different analysis purposes. These insights include: comparison of ratings generated by text with the ratings of the system and the bias issue, extracting important features and their sentiment, querying results based on individual preferences, finding the weakness of the service and so further. The proposed approach is evaluated towards human's logic and it reaches the accuracy 77.8% in feature identification and 70.1% in sentiment score detection. These results are quite comparable with the state of knowledge on feature-based opinion mining.

its is incredibly: are you sure?

I suggest to skip the first paragraph, and just add a sentence that you do research on AirBnb, and in particular compare ratings based on stars and the sentiment expressed in the reviews

TABLE OF CONTENTS

LIST OF FIGURES

INTRODUCTION

The exponential growth of Information and Communication Technologies (ICTs) has played a great role in the development of tourism industry. Electronic tourism (e-tourism) is the application of ICTs for tourism purposes, including the digitalization of all its processes and value chains [?]. Increasingly, ICTs provide users with access to many sources of information and has eventually affected the consumer behavior in tourism industry [?]. In order to remain a strong competitor, service providers have to keep their customers happy. Nowadays there is a huge variety of web services which provide a great complexity and diversity of recommended offers for meeting the demands of travelers. However, the personalized consumption patterns and individualistic lifestyles make it incredibly difficult for the service providers to anticipate tourists' behavior [?]. For decision makers to successfully understand the requirements, needs, desires and preferences of the customer, detailed information has to be obtained. Travelers, making use of the ICT tools that facilitate the information retrieval and decision making processes, have now direct access to all types of information provided by tourism agencies, companies, marketers, enterprises or other users. Furthermore, ICTs and Internet have transformed e-tourism markets from customer-centric to customer-driven [?], meaning that users play a major role in creating and sharing traveling information through blogs and review websites. Online feedback mechanisms, also known as reputation systems, *have emerged as a viable mechanism for fostering cooperation among strangers in such settings* [?]. Examples of these systems, for instance TripAdvisor, Booking.com or AirBnb, after each trade encourage both parties to give feedback about their trading partner based on their own experience. Customer feedback is an essential component in every modern business' tool kit. A big number customer feedback software

Very well done!! You look at the material from different perspectives and introduce it very well.

tools exists for helping service providers to measure and improve customer satisfaction, identify unhappy customers, reduce churn and get valuable insight from customers. However, most of these tools are very expensive, considering that most of service providers can make use of their own feedback mechanisms.

1.1 Problem statement

The most common types of consumer generated feedback are ratings from 1-5 stars and general text comments. An important separation exists between the distinct role of text comments as tacit knowledge and ratings as explicit knowledge and the ways they are analyzed. For online marketplaces to succeed, their feedback technologies must be able to not only collect users feedback, but to properly analyze it and utilize for decision making purposes [?]. However, current online travel systems, aiming to assist the consumer in finding suitable offers, filter the information based on location-price factor and on the overall ratings accumulated from the feedback system, meaning that text reviews are revealed for the public to read but they do not directly affect the overall analysis. Focusing on solely numerical ratings and ignoring the importance of text feedback leads to two major issues for feedback systems.

First, many academic papers on online reputation systems and building trust in the online marketplace report the existence of bias on online reviews [? ? ? ? ?], thus reducing the bias of these systems is an important issue towards a more efficient online feedback system. Utilizing only biased ratings does not necessarily mean that the top result is the most suitable option for a certain user considering the personal requirements. On the other hand, the analysis of ratings does not indicate much information for the service providers, who aim to acquire knowledge on how to improve their services. In order to gain detailed insights from the feedback system, some service providers including Airbnb, ask its users to rate not only the overall quality of the listing, but also six accommodation features. However, since the overall ratings are biased [?], how do we make sure that the ratings for specific features are objective? From this point of view, the bias on quantitative data of the system raises the issue of reliability on the system itself.

Second, the incompleteness of the analysis based solely in quantitative data leads to the need for complementary analysis. Customers' needs are considered multi-dimensional and difficult to measure on discrete scales such as ratings [?], therefore a customer has to extract the needed information from different sources and types of information provided by agencies, companies or other users. According to [?], text comments are particularly interesting for the audience as a new trust-building means in online marketplaces by revealing hidden knowledge, which is often underestimated from their owners and cannot be described by negative/positive ratings. Furthermore, [?] suggests that text opinions influence the decision making process even when the ratings are high. In the Airbnb feedback system a negative rating is followed by a text in 45%

End the sentence here. And tell us in what direction the bias is going, if there is a direction...Then comes the following sentence.

For me bias means systematic error. Thus

of the cases, which implies the great power of text analysis for discovering deeper insights for the listing [?]. Acknowledging the importance of text comments, some feedback systems often offer summaries to all text comments, which mostly consist on a bunch of most used words. However, this bag of words does not necessarily cover the features that a certain user is interested in, neither the features that need to be improved. The users or service providers still have to read all the text comments related to the feature, meaning that it still does not reduce much of the work. A survey by [?] asked the respondents to indicate how many feedback comments they examined before each online transaction. The result showed that 81% reported examining 25 comments (one webpage), 5% viewed 50 comments, 11% more than 50 ones, and only 3% did not examine any text comments. These findings reveal that despite the importance of text feedback to the users, it is difficult for them to access the meaning of numerous text comments [?]. Given this situation, the average human reader will have difficulties on identifying and extracting the relevant information from the opinions in them. Automated analysis systems are thus needed [?].

1.2 Research question

Natural language processing (NLP) enable computers to derive meaning from any human written input, including their opinions. The NLP methods for doing so fall into the category of sentiment mining methods, known also as *opinion mining*. Examples of their application include mainly the movie rating systems (Netflix, IMDb) and the product rating systems (eBay). However, the importance of extracting sentiment of features from comments, besides their overall positive/negative sentiment is often ignored in the literature. This research proposes the implementation of feature-based opinion mining methods from complementing the analysis of customer feedback in e-tourism and accommodation market. The proposed approach uses an ontology based approach combined with sentiment mining techniques for generating opinion scores for each accommodation feature mentioned in the reviews of a feedback system. This solution deals with the two issues mentioned above, bias of ratings and the need for complementary text analysis for feature extraction. From its point of view, both issues can be brought together as one, since text analysis is believed to contribute to a better rating systems by reducing its bias [?].

This paper answers first the question of

how can opinion mining methods be aligned with the quantitative data analysis in order to enhance focus on customer feedback and produce detailed analysis results.

By estimating sentiment scores for the text reviews, the pipeline transforms the text data into discrete quantitative form, which can easily be analyzed for different purposes. Some of the implications of the analysis are treated in the next sections. In addition to this analysis, it is important to find out *how good can feature-based opinion mining estimate the quality of accommodation features in e-tourism feedback systems*, which measures the reliability of the

"for" I
guess

put this
as a
quote
too

analysis. By answering these two questions, the purpose of the paper is to offer to service providers a new reliable approach on enhancing focus on customers, based on users' generated content.

1.3 Research rationale and structure

Customer Focus Theory is one of the essential parts of Information Studies (IS) and it serves as a guidance for business on how to put focus on their customers, as their most valuable asset. From Figure ?? can be clearly seen that Customer Requirement, Information, Feedback and Relationships are the four key factors of a good focus on customer [?]. In alignment with this guidance, businesses try their best on gathering customers information and feedback in several ways starting online forms, surveys and interviews to social media, forums and blogs. However, the most important part of work comes exactly after having the feedback, explicitly on how it is analyzed, interpreted and utilized. Most of business pay huge amount of money on outsourcing the feedback analysis, human resources (like the Amazon Mechanical Turk) and expensive software. From this point of view, the proposed approach is an added value for both service providers and their customer for several reasons. First, an analytic tool taking care of text feedback and bringing it into the traditional quantitative type of data, saves both a lot of time and money. Second, the insights and analysis are personalized to the business purposes, meaning that the winner is the one who gets the requirements properly. Thirdly, transforming text feedback into quantitative data serves as a measurable data-drive target over time, where businesses can check the quality directions of their services.

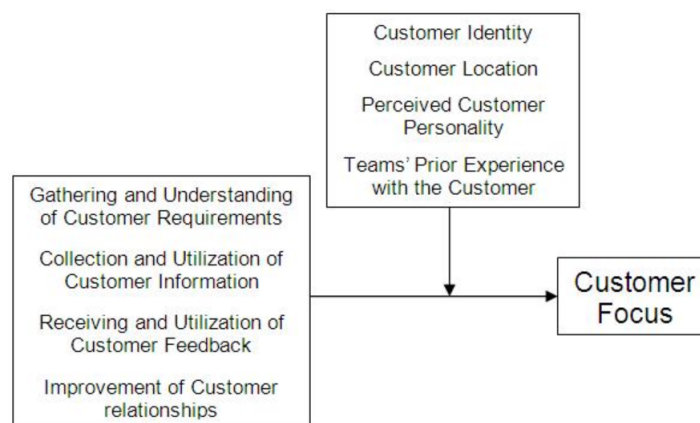


Figure 1.1: Customer Focus Theory

Put the source of this picture in the caption using a citation

From the literature point of view, it is noticed a gap in the alignment between text analysis with the analysis from quantitative data. Research in this field would first add value to the Customer Focus Theory and secondly to the use of feature-based opinion mining for analysis purposes. The existing systems and methods of feature-based opinion are examined and presented

in the following chapter. Shortly, we can conclude that the literature lacks a system where all the features belonging to a certain domain are analyzed and scored discretely based on text feedback. Furthermore, [?] suggests the use of sentiment analysis for complementing quantitative data analysis as future work to be done, in order to reduce the bias of feedback systems. This proposed approach examines this suggestion and can be used for further similar hypothesis.

To provide an answer to the research questions, this paper is organized in seven chapters. The importance of gathering and analyzing customer feedback is introduced by the first part of the paper and it is followed, in the second chapter, by the current state-of-the-art of feature-based opinion mining used for analyzing customer feedback. The proposed approach is discussed in the third chapter by explaining each step of the pipeline developed for this research. The fourth chapter covers the methodology used in the research, from collecting the data to the analysis of the output. In chapter five, the whole proposed approach is evaluated in comparison to human logic. Afterwards the results are presented and visualized, which leads the discussion to the limitations of this work, its implications and further work to be done. Finally the paper concludes with a recap of the main findings.

What kind of paper is this? Do you have it? Are these your friends? I suggest you explain this, as you mention it quite often

CURRENT STATE OF KNOWLEDGE

The task of extracting opinions from customer reviews is an issue, which has the attention of research communities for more than one and a half decade. With the continuous growth of Internet use, numerous social networks, blogs, news reports, forums, e-commercial websites and other platforms serve as an open-space for expressing opinions. The domain of interest varies from general public events to political campaigns as from product choices to marketing strategies. Sentiment analysis (SA) is the study of opinions, emotions and attitude towards an entity expressed in free text [?]. Sentiment analysis itself is a problem that includes several inter-related aspects, such as subjectivity classification, polarity determination, review usefulness measurement, opinion spam detection and so further. Due to the lack of standards in the methods used and the complexity of the problem, it is difficult to portray and categorize all the views presented in the literature. From 162 research articles covered in [?] during the time frame 2002-2015, my main focus are the feature-based opinion mining systems.

The reviewed systems base their extraction on statistical methods, unsupervised learning methods or ontologies. Statistical methods typically apply some rules on the frequency of nouns to find explicit features, unsupervised learning method group the explicit features and the similarities by measuring distributional properties of words and ontologies provide a pre-set list of features for the systems to be based. On the other side, there are mainly two approaches to polarity determination: machine learning and lexicon based methods. The level of sentiment analysis here differs from document level, sentence level or word level. The examined systems vary from each other on the methods used for feature identification, the granularity of sentiment analysis, the methods used for polarity detection or the way the output is presented. Due to the fact that this paper makes use of an ontology based approach, the examined systems are grouped

into non-ontology based and ontology based, although it is hard to compare the performance of these systems due to the lack of standard test data and methods. Starting from non-ontology based approaches, Liu and Hu introduced in 2005 Opinion Observer, a lexicon-based system for sentiment analysis of opinions [? ?]. Their system extracts product features by using association rule mining, explicitly the Classification Based on Association (CBA) algorithm. The system only uses adjectives as opinion words and assigns their polarity based on WordNet dictionary. The polarity of an opinion expression is determined in sentence level based on the number of positive/negative words. The final output are stored in a database in the form of (*feature, number of positive expressions, number of negative expressions*). The system reaches 90% precision and 83% on recall.

In 2007 two other systems are proposed, namely Red Opal [?] and the famous OPINE [?]. Red Opal aims to score product features based on customer reviews. It uses probability-based heuristics for identifying frequent nouns and noun phrases for feature extraction and it is the only systems which makes use of review stars for defining the sentiment of features. In other words, the rating a reviewer would give are not calculated as an overall score but are assigned to specific features mentioned in their own reviews. A user interface is used for offering to the customers a ranked list of products based on the selected features. The precision of Red Opal is 88% and the recall is not measured.

The second approach, OPINE [?] is based on KnowItAll information extraction system which is domain-independent. This approach extracts explicit product features using the Pointwise Mutual Information (PMI) between phrases. It uses explicit features to identify potential opinion phrases based on assumption that "an opinion phrase associated with a product feature will occur in its vicinity on syntactic parse tree" [?]. After the extraction of the opinion expression, the semantic orientation of opinion words is determined by using relaxation labeling, an unsupervised classification technique. The result is a set of *feature, list of opinion sentences*. The precision of OPINE is 79% and the recall 76%.

Weakness Finder [?] is one of the well-known proposal of 2012. It aims to help the identification of features, and group these features into different aspects by using explicit and implicit features grouping methods, then judge the polarity of each sentence by using sentence-level sentiment analysis. The methods used for feature extraction are both collocation statistic based methods and lexicon-based on HowNet dictionary, thus a hybrid approach. The sentiment of features is detected in sentence level as positive, negative or neutral extending Liu's work, Opinion Observer system [?]. Weakness Finder has a precision of 85.6% and recall 85.2%.

Bagheri et al. proposed in 2013 a novel unsupervised and domain-independent model for detecting explicit and implicit aspects in reviews for sentiment analysis [?]. In the model, first a generalized method is proposed to learn multi-word aspects and then a set of heuristic rules is employed to take into account the influence of an opinion word on detecting the aspect. Second a new metric based on PMI and aspect frequency is proposed to score aspects with a

please
indicate
on what
task

You may
also cite
the work
of me
and my
room-
mate
Prof
Mokken,
and
also our
paper
"Words
with at-
titude":
go to
Google
scholar
and look
at my
most
cited
works

new bootstrapping iterative algorithm, works with an unsupervised seed set. Utilizing extracted polarity lexicon, the approach maps each opinion word in the lexicon to the set of pre-extracted explicit aspects with a co-occurrence metric. The output would look like in the example (*great* - *phone, price*); (*good* - *battery life, sound quality*). The proposed model was tested on a set of English product review and it reaches precision 87.5% but low recall 65%.

The first proper system using an ontology-based approach is proposed in 2012. Fuzzy domain ontology sentiment tree (FDOST) [?] aims to discover sentiment polarities over product features based on ontologies, which unites traditional natural language processing (NLP) techniques with sentiment analysis processes and Semantic Web technologies. The authors first construct the hierarchical relationships among product's attributes, product's attributes and corresponding sentiments through fuzzy sets. Then they extend the hierarchy of FDSOT, including the extraction of sentiment words, product features and the relations among features. Finally, the polarity weights of features are assigned making use of Double Propagation. The output of the FDOST are the features and the polarity weights, for example *{cheap, +1626.729}* reveals that there are 1626 consumers expressing positive opinion about the price.

In 2014 Penalver et al. proposed an another "novel" ontology based methodology for feature extraction [?]. Their system has four stages: a) feature identification based on ontology mechanism; b) polarity assignment to each feature based on SentiWordNet and the relative position in each user's opinion; and c) a new approach for opinion mining based on vector analysis. The resulting polarity value is an Euclidean vector with three coordinates (x, y, z) for each feature counting the number of positive, negative and neutral opinion respectively. The sum of these vectors will be the global polarity expressed by a user or for a certain feature. Their approach has achieved an accuracy of 89.6% , however the sentiment classification of all the opinions is based on document level, meaning that for each feature at least 100 sentences are needed to extract its sentiment.

Lastly, in 2015 is introduced Type-2 fuzzy ontology called T2FOBOMIE, a novel extraction and opinion mining system based [?]. The system reads the customers' full-text query for hotel search and reformulates it for extracting the user requirement into the format of a proper classical full-text search engine query. The proposed system retrieves targeted hotel reviews and extracts feature opinions from reviews using a fuzzy domain ontology. T2FOBOMIE uses a classical extensible markup language (XML)-based ontology, Protege OWL. The sentiment of features is identified in lexicon-based method, making use of SentiWordNet for detecting sentiment score and then grouping it into *positive, positive, neutral, negative, negative*. The results are presented with the help of an user interface which requires the customer to write his full-text query and the output of the system would be a list of hotels, the corresponding sentiment polarity and a link to the website. The system reaches the precision 80.8% and recall 60.3%.

Table ?? gives an overview of the comparison within these systems based on methods they use for feature extraction, sentiment detection, domain of usage, language and output format.

SYSTEM	FEATURE EXTRACTION	SENTIMENT DETECTION	DOMAIN	LANG	OUTPUT	P/R
Opinion Observer (2005)	Association miner, CBA	Lexicon-based (WordNet)	Product	English	Feature, # of positive expression, # of negative expression	90% 83%
Red Opal (2007)	Probability-based heuristics	Assign star rating	Product	English	User interface of ranked results	88%
OPINE (2007)	Unsupervised, Web PMI	Relaxation labeling	Ind.	English	List of sentiment sentences	79% 76%
Bagheri et al. (2013)	PMI, DBA	Bootstrapping	Product	English	A co-occurrence matrix: (feature; opinion words)	87.5% 65%
Weakness Finder (2012)	Collocation statistics	Lexicon-based (HowNet)	Product	Chinese	List of features with negative sentiment	85.6% 85.2%
FDSOT (2012)	Fuzzy set, Ontology-based	Double Propagation	Product	Chinese	Set feature;polarity	72.4% 62.9%
Penalver et al. (2014)	Ontology-based	Dictionary-based (SentiWordNet)	Movie	English, Spanish	Euclidean vectors of polarity	89.6% ¹
T2FOBOMIE (2015)	Ontology-based	Lexicon-based (SentiWordNet)	Hotel	English	List of hotels with positive/negative polarity on feature	80.8% 60.3%
The proposed approach	Ontology-based & lexicon-based (WordNet)	Rule-based and lexicon based (VADER)	Hotel	English	Matrix of discrete scores per each sentence and features	77.8% 79.2%

Table 2.1: Comparison of examined systems on feature-based opinion mining

Most of these systems, no matter if they are ontology-based or not, have high precision and recall, however it is noticed that the recall is usually lower than precision.

These approaches are mostly found for English or Chinese reviews. A big limitation of the systems based on statistical and unsupervised learning methods is that they ignore features with low frequency of occurrence. This fact is not anymore an issue for the ontology-based systems, however they vary on the methods used for building the ontology as there lacks any unified, standard, multi-domain ontology. The sources used for building the ontology of my proposed approach will be described in the Methodology chapter. Secondly, the examined systems detect the sentiment of features based on positive, negative and neutral polarity and not in discrete scores. Grouping the polarity in this form guarantees better scores in the accuracy of algorithms, however it does not leave room for further analysis with the data. In addition to this, the third limitation is the way of assigning sentiment to features. Depending on sentence or document level, an overview of the number of positive/negative opinions, the sentiment sentences itself or either their summarization, do not contribute in getting deeper insights from text, neither reduce the workload of the customers and service providers for analyzing the data. Thus, this paper proposes the use of estimated discrete scores of sentiment per feature in order to get deeper insights in customers opinion from free text and align this with the quantitative data retrieved from the feedback system. This proposal will be explained in the following chapter.

I like your table very much, and it gives a good overview. But I am afraid I do not understand what you mean by precision and recall. I am also afraid that it

THE PROPOSED APPROACH

Good setup, but make it less sloppy. Here you want to be very precise, as you want that other people can easily redo your work

For transforming the text feedback into quantitative data, which will serve as input for analysis purposes, this paper proposes an approach as shown in Figure ?? . The pipeline consists of four main steps: *pre-processing*, *feature identification*, *sentiment detection* and *data analysis*. Initially, the whole input of the pipeline is the huge corpus of text reviews stored in Neo4J database. This big amount of data from social is considered to be very noisy, therefore it will be cleaned up as be described in the following section. Features and sentiment are then identified in text based in sentence level. The output of these three stages consists of the quantitative sentiment scores for sentences of reviews and the features identified in them. This output will then serve as the input data of the final stage of the pipeline: data analysis. The whole pipeline is built using Python programming language and its related packages. All the data analysis and visualization are written in Jupyter notebook. The following sections will explain each step of the proposed approach in further details.

3.1 Pre-processing

The pipeline reads the text data from cloud with the help of *py2neo*, a toolkit for working with Neo4J from within Python applications, and it formulates the queries in Cypher, the querying language for Neo4J graph database. The reviews in the corpus are read one by one and each of them is checked if it fulfills the language requirements. In my work, the whole concept of the pipeline is developed in English, as the most used language in e-tourism websites and as the easiest language for text mining. Thus, every time that the algorithm runs into a review

using which software and add a reference to the ena-

not in English, it will ignore the review and continue with the next one. For each English review detected, the algorithm will cut the text into sentences , with the intention of detecting

How?
this is
a non
trivial
task

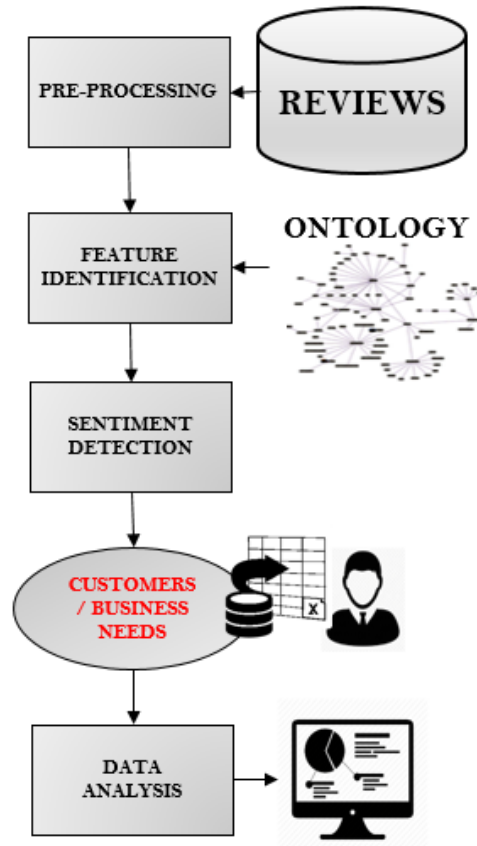


Figure 3.1: The proposed approach

the sentiment in sentence-level and extracting features within these sentences. Afterwards tokenization and lemmatization are performed to each sentence of the review. Tokenization is used form chopping the sentences into words, phrases, symbols and emoticons, each represented as a token. Lemmatization is the process of getting the root of the words, called lemmas, while ignoring its other forms as part of the speech. For tokenization the pipeline uses the *TweetTokenizer*¹ package, part of NLTK and Twitter-aware designed, in order to be able to identify emoticons and adapt to new domains. On the other hand lemmatization is performed using *WordNet*², a large lexical database for English terms, in which nouns, verbs, adjectives and adverbs are grouped into sets of cognitive synonyms (synsets), each expressing a distinct concept [?]. The output of these two steps is a list of lemmas for each sentence, which will be used by the next step in the pipeline as explained in the following section.

¹<http://www.nltk.org/api/nltk.tokenize.html>

²<https://wordnet.princeton.edu/>

3.2 Feature identification

As soon as each sentence in the corpus of reviews is represented as a list of lemmas, an ontology based approach is used to identify the accommodation features of each sentence. The ontology is chosen for feature identification as a way for detecting all the terms, concepts and relations linked to accommodation domain. Prior to identifying the features, the pipeline builds a list of synonyms, hyponyms and hypernyms for every term of the ontology, in order to increase the range of matches found in the text. The list is based on the Synsets and relations between concepts as introduced by WordNet library. The lists of related terms is lemmatized, as explained above, and the duplicates are removed. The lemmatization process aims to create a list of related lemmas for each feature in the ontology, which would then be used for feature identification in the sentences. Thus, the feature identification steps is a string match of the list of lemmas for every single sentence, with the list of related lemmas of every feature in the ontology. Considering that the ontology consists of many terms, this sample pipeline is trained to identify only six features that Airbnb asks the customers to manually rate as part of their feedback, namely *accuracy*, *cleanliness*, *check-in*, *communication*, *location* and *value*. However, the conclusion are drawn equally as all the features of ontology are included. When a word in the sentence is identified to be a feature, the algorithm jumps to the next word of the sentence. Within one sentence more than one feature can be identified, as well as features can be mentioned more than once. Therefore, for every possible feature match, a *match counter* variable which keeps track of the features mentioned is assigned to the sentence. Here, the pipeline considers each sentence as independent and it ignores the logical connection between two sentences of the same review.

3.3 Sentiment detection

A very important part of the pipeline is sentiment detection. The algorithm used for this purpose is VADER (Valence Aware Dictionary for Sentiment Reasoning) ³, part of Python packages. VADER is a lexicon and rule-based sentiment analysis tool that is specifically attuned to sentiments expressed in social media. For building the valence scores for sentiment intensity, VADER considers several well-known sources, such as ANEW, SentiWordNet and SenticNet. These dictionaries include not only lexical features, but also grammatical and syntactical rules. Therefore, VADER is able to deal with negation, capital letters, degree modifiers, emoticons, punctuation types, contrastive conjunctions (*but*, *however* and slangs. Based on the comparisons of 22 sentiment mining tools of the last decade, VADER is ranked as the best algorithm for comments and the second best for social networks [?]. Some of its assigned scores are (*good* 0.7); (*great* 0.9434); (*awesome* 0.8306); (*dirty* -0.83066); (*terrible* -0.9434) ⁴. In this paper, the pipeline uses VADER to return a sentiment score for every sentence of the review, which would afterwards

³<https://pypi.python.org/pypi/vaderSentiment>

⁴The full lexicon can be found in: https://github.com/cjhutto/vaderSentiment/tree/master/additional_resources

Here you talk too much on the implementation while I like to see a formula which exactly explains the feature scores you give to each sentence. As your thesis did not have

serve as a discrete score for the identified features in that sentence. Considering that in one sentence, more than 1 feature can be identified or the same feature can be identified more than once as mentioned above, a probabilistic model is developed. This model serves for defining the probabilities that a sentiment score would reflect user's opinion on each explicit feature identified within one sentence.

The dataset consists of 2 356 listings, $L = \{L_0, L_1, L_2 \dots L_{2355}\}$. Each listing in this corpus can be represented by the ID of the listing in the database lis_id . Every listing from this set has a number of reviews, which varies from one listing to the other and is identified from ID of the review $R_{lis_id} = \{R_0, R_1, R_2 \dots R_r\}$. For a single review of a certain listing would be R_{lis_id, rev_id} . This review is composed by a number of sentences $S_{lis_id, rev_id} = \{S_0, S_1, S_2 \dots S_s\}$. The pipeline is trained to detect the sentiment of all the sentences, which will be in the same form as the one above. In each of these sentences, the pipeline is trained to identify six accommodation features $F = \{accuracy, check-in, cleanliness, communication, location, value\}$. When the sentiment of a sentence is detected, it does not particularly refer to a certain feature. In order to calculate the part of the sentiment that belongs to a certain feature identified in the sentence, we use a probabilistic model. According to this model, the probability that the compound sentiment score refers to a certain feature is p_i , where i is the index of the feature from the list F . This probability is uniformly distributed between the identified features and is calculated as $1/k$, where k is the number of features mentioned. Likewise, if the feature is not identified in the sentence the probability would be 0.

$$Sentiment = \sum_{i=1}^n p_i * sentiment_i \quad (1)$$

This formula for calculating the sentiment of the features is part of the sentiment detection phase of the pipeline and it is repeated for every single sentence. For instance, considering the reviews of listing 24328, $R_{24328, 10146}$ would represent the review: "We had a great stay at Joe's. The handy guide to the house and neighborhood was much appreciated, as were Joe's clear instructions for checking in while he was away. The house was comfortable and eclectic, full of personal character." Thus, $S_{24328, 10146, 1}$ represents the second sentence of the text above. In this sentence can be identified three features *check-in*, *communication*, and *location*. The compound sentiment score of the sentence above has in this way to refer to one of these features in a probability of 33.3%. Figure ?? shows the results of the pipeline for this case.

3.4 Data analysis

The last step of the pipeline and the most interesting one for service providers and businesses is data analysis. Up to here, we saw the pipeline processing all the text data step by step and transform it into meaningful sentiment scores for each sentence and each accommodation feature

This is a very rough cut from one part to another

I propose you put some information on the dataset at the beginning of this section. Where does it come from? From which town(s), who collected it? What are the median, mean, max number of reviews per listing? And

SENTENCE	SENTIMENT	ACCURACY	CHECK-IN	CLEANLINESS	COMMUNICATION	LOCATION	VALUE
The handy guide to the house and neighborhood was much appreciated, as were Joe's clear instructions for checking in while he was away.	0.851896	0	0.28394	0	0.283936943	0.28394	0

Figure 3.2: Example of the probabilistic sentiment results

of the reviews. The output of these phases is a Comma Separated Vector (.csv) file, consisting of all listings, all listings' reviews, its respective sentences and the sentiment for the identified features on these sentences. These scores are analyzed with Pandas⁵, an open source library providing high-performance, easy-to-use data structures and data analysis tools for the Python. The workspace of analysis and visualization of results is Jupyter Notebook, a web application that allows users to create and share documents that contain interactive code, equations, visualizations and explanatory text. The data analysis in this step includes calculation and visualization of sentiment scores per reviews and listings, frequency of ontology-based features, ranking of listings based on overall sentiment or based on sentiment of one or more specific features, ranking of features for one listing based on their sentiment, identification of listings with the biggest number of reviews and identification on listings where the host have canceled once or more the reservation, computation and comparisons of sentiment scores from the pipeline with the ratings of Airbnb and so further. An example of the analysis, Figure ?? shows all the listings with low rating (less than 4 stars since the lowest rating is 3 stars) and visualizes the differences between the sentiment scores generated from the pipeline and the ratings found on Airbnb. The analysis can of course be extended and customized to match the specific interest of the service providers or the customers. The main results of this analysis will be explored in the results section, followed by the implementation of pipeline according to its usefulness for service providers and customers respectively.

Really
you
have
not more
than 20
listings
with
such low
score ,
so less
than 1%

⁵<http://pandas.pydata.org/>

```
In [31]: # Compare the ratings and sentiment scores from pipeline for listings which have less than 4 stars
low_rating = comparison[comparison['Sentimentscore']<4]
low_rating[['Overall','Sentimentscore']].plot(kind='bar')

Out[31]: <matplotlib.axes._subplots.AxesSubplot at 0x761dd50>
```




Figure 3.3: Example of data analysis in Jupyter Notebook with IPython & Pandas Library

MM: did not look at this section yet

The proposed pipeline for opinion mining and extracting features from text feedback is written in Python 2.7. In addition to basic packages, I have made use of NLTK 3.2.1, py2neo, vaderSentiment, ipython, jupyter, langdetect, pandas, matlab etc. All the codes belonging to the four phases of the pipeline, including data analysis and visualization can be found in Appendix (add number). Besides what is already explained in the proposed approach chapter, the next section describes two main data sources used in the pipeline: the dataset from Airbnb and the accommodation ontology.

4.1 Dataset

The algorithm deals with a huge set of reviews, which are retrieved from the feedback system of Airbnb and they serve as the input of the pipeline. The whole corpus of reviews consists of 3.4GB, including reviews of the listings from the Netherlands and United Kingdom. This corpus is then filtered to be focused only in the data of Amsterdam, for simplicity purposes. The same code can run in the bigger corpus, with the help of a more powerful machine. The data set is stored in cloud, in the Neo4J graph database, where the nodes can represent a listing (accommodation offer), a guest/host, a review or a response.

4.2 Accommodation ontology

In knowledge management and Semantic Web research areas, ontologies are considered essential in order to describe various concepts and relationships between them. Ontologies are formal specifications of a shared conceptualization of a domain. For the accommodation domain and related domains a few ontologies have been proposed, which cover different aspects in this domain. However these ontologies are often found as sub-ontologies of the e-tourism domain. In this research, the ontology of accommodation is based on several sources. Firstly HONTOLOGY [?], which is brought in alignment with Accommodation in QALL-ME, Tourist Accommodations in DBpedia.org and Lodging Business in Schema.org, was initially created by following different scenarios of booking an accommodation and processing reviews in Web services. Secondly ACCO Accommodation Ontology [?], which is an extension of GoodRelations ontology [?], is based on Owl Ontology Language (OWL) and is supported by Google and Yahoo for the e-commerce accommodation offers. Besides these ontologies, further covers features mentioned in the Web based accommodation services of AirBnb.com, Booking.com and TripAdvisor.com.

This needs a lot of cleaning up and making things precise.

ratings given by humans sounds better to me. Also change this in the introduction. Logic is something else

or 100 sentences?

For checking the efficiency of the proposed pipeline, its performance is compared to human logic for two main purposes. First, the evaluation aims to find out how well VADER, the selected sentiment detection algorithm, performs in the chosen dataset and secondly, how well can the pipeline identify the features of a certain sentence of review. To answer these two questions, a sample of 100 text reviews are randomly chosen from the dataset and are given to humans for evaluation. Each respondent is required to first read the sentences, then to estimate a score of sentiment for each of them on a scale from -1 to +1 and finally to mark the accommodation features, which the sentence in question refers to. These tasks are similar to what the pipeline is programmed to do. The evaluation form is completed by 5 humans, who come from different educational backgrounds, are geographically diverse and are Airbnb users. Their answers are analyzed and compared to the pipeline output using SPSS. Considering the diversity between users and in order to understand the common human logic, their evaluations are firstly compared to each other using correlation values between samples and distribution statistics. This analysis showed that three of the respondents had very similar answers, consisting in a correlation varying between 72% to 83 % between each other and a very similar distribution curve. Based on these results, the chosen sample to represent the average human logic are the average rates of the three "unbiased" respondents.

5.1 Evaluation of VADER algorithm

VADER is considered to be one of the best algorithms for sentiment mining with a very high

accuracy [?]. However, when VADER was tested to perform in 22 different datasets, it was noticed an overall accuracy of 78.7%, which varies significantly from one dataset to the other [?]. To clear all the doubts of how VADER would perform on reviews and specifically in the chosen dataset of this research, its scores are compared to the average human rating. The correlation between the two sets based on cases results to be 70.8% . Although, this value is lower than the average case of VADER's accuracy , it indicates a satisfactory scale of accuracy for the Airbnb dataset. From Figure ?? can be noticed that the two datasets share the same mean (4.125 compared to 4.13) and have similar standard deviations.

The differences between VADER scores and human evaluation are also calculated and they




Figure 5.1: Distribution of VADER and humans' sentiment scores

result in 53% of cases when VADER has detected the exact score as human logic, 24% cases where the difference is just half a star, 14% cases with one star difference and the 9% cases with more than one star difference. The last ones form the group of "errors" in VADER sentiment scores. The highest value of error is 2.5 stars, in only 3% cases, which means that VADER will never consider a very positive sentence as a very negative one and the other way round. However, it may consider these kind of sentences as neutrals or may be confused of the sign of sentences with slight sentiment.

5.2 Evaluation of feature identification

The effectiveness of the proposed technique for feature identification is measured by using precision, recall and accuracy as suggested by [?]. The values are calculated for each of the features as:

$$Precision(p) = \frac{TP}{TP + FP} \quad (2)$$

this sentence is completely unclear. You have 4 measures of sentiment for 100 sentences. What did you do? Explain it EX-ACTLY.

You compare accuracy with correlation. ?????

They also have the same median: mention that and its value

I think you should calculate the

$$Recall(r) = \frac{TP}{TP + FN} \quad (3)$$

$$Accuracy(a) = \frac{TP + TN}{TP + TN + FP + FN} \quad (4)$$

where TP (true positives) is the number of sentences that the algorithm correctly identifies the right features; FP (false positives) is the number of sentences that the algorithm falsely extract wrong features; FN (false negatives) is the number of sentences that the algorithm fails to identify the right features and TN (true negative) is the number of sentences that the algorithm correctly does not identify any feature. The average human logic identifies a feature in a sentence when it is agreed by the majority of respondents. In other words, *accuracy* would refer to the portion of correctly predicted sentences against all sentences for evaluation. *Precision* would show how many sentence identify a feature where it is not mentioned, and recall shows the portion of sentences that the pipeline fails to identify the right feature. Ideally the algorithm shall have values of precision and recall close to 1 for each of the features. This evaluation of the pipeline as illustrated in Figure ?? shows that the algorithm manages to identify some features better than others. For example the *value* and *cleanliness* features of the listing are identified almost always with high precision and recall, but a low precision for *communication* means that the algorithm retrieves many FP. The opposite happens for *check in* when the precision is very high but many sentences fail to be identified. For the six features overall the algorithm reaches the accuracy and precision 77.8% and recall 79.2%. Future work needs to be done in improving this step of the pipeline.




Figure 5.2: Precision/Recall matrix for feature identification

Are you sure about this? So if a sentence is about A,B and C, and the algorithm only says A and B, it is completely wrong? This sounds weird! You must count whether you have correctly spotted a feature. Change all this

RESULTS

Sentiment analysis and feature extraction from text reviews of Airbnb feedback system, provides us a set of discrete estimators for customers' opinions. The sentiment scores can vary from -1 (extremely negative) to 1 (extremely positive). From the analysis of reviews on Amsterdam dataset the most negative sentiment is -0.964284 and the highest score 0.998195. In total there are 12 798 negative sentences (sentiment lower than 0) and 236 904 positive ones, meaning the number of positive sentences is almost 19 times higher. This can be clearly noticed in Figure ?? by comparing the density of positive and negative scores. In addition, we can also notice a tendency of sentiment to avoid the "slightly" negative or positive zones around 0. The sentiment




Figure 6.1: Sentiment scores of each sentence of reviews in the dataset

scores are also grouped by accommodation features and for each of them its seen the distribution

of these scores. By comparing these distributions to each other and to the overall sentiment scores, it is noticed that the more the feature is mentioned the more similar the distribution is with the overall sentiment pattern. However, small differences are recognized for each feature, for example the values of feature *cleanliness* have a tendency to be mostly positive and grouped between 0.3-0.5 or 0.7-0.9. These patterns of distribution can all be found in Appendix I.

An interesting point of view is the analysis of what features are most mentioned in proportion with the number of comments in the listing. This analysis is done in three levels: sentence level, review level and listing level. For all the three levels, we can clearly see that *location* is the most mentioned feature by reviewers of Airbnb and *check-in* is the least mentioned one. From this comparison we can also indicate that the frequency of mentioning features is more accurate in review level, as in the first case the number of sentences is very high which makes the percentage be under-estimated. On the other hand, in listing level it would take at least one sentence to indicate that the feature is mentioned in the listing, no matter if at one listing is mentioned 100 times, which can cause over-estimation. Following this logic, Figure ?? represents for each listing




Figure 6.2: Percentage of sentences, reviews and listings where the features are mentioned

the number of reviews which mention the features and it compares it to the overall number of reviews. In this way, we would prefer to choose listings where we have many reviews for a certain features i.e. *location* compared to the overall number of reviews in order for the sentiment score to be more reliable. For each of the listings the compound sentiment scores over a certain feature are easily calculated. The second part of Figure ?? shows for the same ten listings the compound sentiment scores of each feature and the compound sentiment. In this way, the sentiment scores can be compared with the number of reviews. From the Figure we can see that for the listing with ID 22886, even though it has the highest sentiment score on *communication*, it is reviewed by only a very few people, compared to the other listings. This result means that the sentiment score of the first listing would be more reliable than the one with a few reviewers but a good sentiment score. In this case each customer would have to make an individual choice for which listing they would go by considering both the number of reviews on a certain feature and their sentiment.

Another aspect of analysis from the generated sentiment scores is their comparison with the Airbnb quantitative data. In the Airbnb feedback system, the customers can see the number




Figure 6.3: Features mentioned in review level per each listing and their sentiment

of reviews for a listing, the overall average rating after the first three reviews and the average rating of the six specific features: *accuracy*, *check-in*, *cleanliness*, *communication*, *location*, *value*. These ratings are compared to the ratings generated by the pipeline by converting the sentiment scores $[-1;1]$ in the scale 1 to 5 stars. From the analysis we can notice that the overall ratings of generated from the pipeline are quite similar to the ratings in the Airbnb system.

It is obvious from Figure ?? a) that there is no listing with negative sentiment, lower than 3 stars, both in Airbnb system and the pipeline results. Furthermore there is a tendency of Airbnb ratings to be higher than the ratings generated from the text feedback, concretely in 89.7% of the cases, consisting in a difference of half a star in 50.6% of cases, one star in 38.7% of cases and only 0.3% of difference more than one star. The standard error of the mean (SEM) between the pipeline sample and Airbnb is 0.0096 which indicates a precise estimation of compound ratings. These results prove the precision of the pipeline in predicting the overall rating of the listings based on text feedback and in the same time indicate a factor of reducing bias in the Airbnb

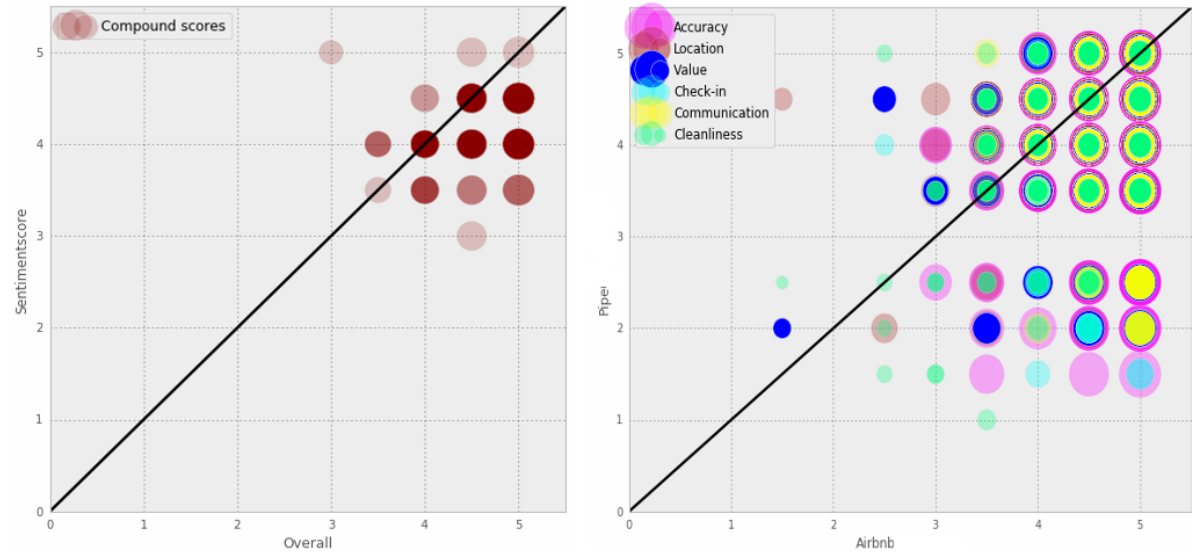


Figure 6.4: Comparison of overall and features rating distribution

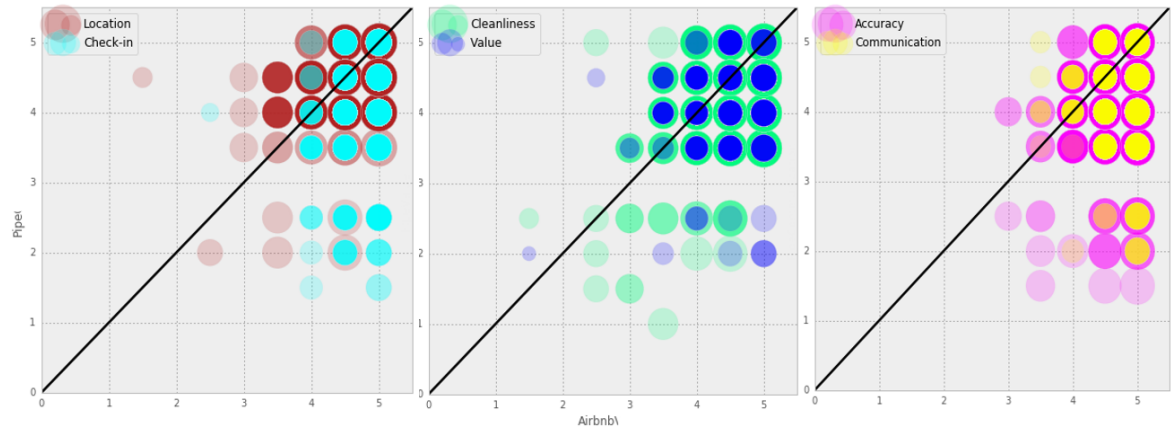


Figure 6.5: Features rating distribution detailed

system, as suggested by [?]. The only outlier for the compound ratings of the pipeline compared to Airbnb is a case where the sentiment predicted for a listing is 5 stars, when the Airbnb system assigns only 3. Since this case is isolated and occurs only once, we have considered it in the range of errors of the pipeline precision.

Similarly in feature level we can also see that most of the rating values of Airbnb and pipeline fall in the upper corner of the quadrant, however it is obvious the presence of a big number of outliers. The ones that interest us are the values for which the pipeline estimates negative ratings (lower than 3) but in Airbnb those features are rated with the maximum number of stars. From Figure ?? b) and its detailed version Figure ??, we can notice that this phenomenon is present at every feature, but most likely for *Check-in*, *Communication* and *Accuracy*. The findings are also supported by the standard error of the mean (SEM), which is higher for these

three features and the lowest for *location* and *communication* respectively 0.01 and 0.008. These findings align with the fact that *location* and *cleanliness* resulted to be the most mentioned features in the reviews, therefore the higher the number of cases in the sample the bigger the precision. Furthermore, from the evaluation of the pipeline we saw that feature extraction is most accurate and precise for *cleanliness* therefore the results are more reliable. On the other hand, for *check-in* the evaluation of pipeline suggests that half of the cases mentioning this feature are missed by the pipeline. Therefore, the explanation of the outliers can lay between the accuracy of the pipeline on extracting features from the text and the real frequency of features mentioned in reviews. Thus, the better the feature identification from pipeline and the higher the feature is mentioned in reviews, the smaller is the number of outliers in distribution of pipeline generated ratings per feature compared to Airbnb and overall ratings.

Another aspect that the analysis of text feedback can be used is for checking the number of times when the host has canceled the reservation. Every time that the reservation is canceled an automatic message is posted in the text feedback space as a review by the guest-to-be. The occurrence of these types of messages can be check for all the listings and from the analysis we could retrieve the listings which have the highest probability of canceling. It was noticed that in a rare case the number of cancellations reached 27, however the mean number of cancellation would be 1.7 per listing. In total there are 32.8% of the listings which have canceled at least once the reservation, meaning that the customers would have to pay attention to these cases in order to let their plans spoiled.

Many other insights can be drawn from the analysis of the sentiment scores and features extracted, which are similar to the analysis of quantitative data as we deal only with discrete sentiment scores. These results showed the alignment between the quantitative data of Airbnb as star ratings with the output of the pipeline. Furthermore, they showed that the text feedback can reveal hidden knowledge which can not be obtained only by analyzing the ratings in the feedback system. The analysis involves both the customer and business perspective therefore it is seen as an interesting aspect to explore further than this research.

Here will go the discussion.

- How my work is similar to others and where it differs
- Recap of the results and their weakness
- Recap of evaluation

7.1 Limitations

- Language only english
- Could have used more algorithms and compare them
- More people for evaluation
- Accommodation ontology and query search n-gram
-

7.2 Implications

The main contribution of this work is threefold. First, an ontology- based feature identification permits the reuse of existing vocabularies and ontologies in order to extract feature related information from opinions in different domains. Second, four different configurable methods

for feature polarity identification are proposed. These methods can be configured with different parameters to obtain the best polarity identification approach for different domains and languages. Finally, the vector analysis based opinion mining approach permits the sentiment classification of a document to be calculated, and no training phases are necessary.

7.3 Future work

- The beginning of a recommendation system
- Extend the algorithm to n-grams
-

Here will go the conclusions.

APPENDIX A

Begins an appendix